

# Kinney\_DSC550\_Final

January 26, 2020

## 0.1 D. Kinney DSC 550 Final Project

### 0.1.1 Part 1: Graph Analysis

---

```
[1]: import warnings
warnings.filterwarnings("ignore")

[2]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import sklearn.linear_model
import statsmodels.formula.api as smf

from pandas.plotting import scatter_matrix
from plotnine import *

pd.set_option('display.max_columns', None)

%matplotlib inline

[3]: def prepare_country_stats(oecd_bli, gdp_per_capita):
    oecd_bli = oecd_bli[oecd_bli["INEQUALITY"]=="TOT"]
    oecd_bli = oecd_bli.pivot(index="Country", columns="Indicator",
    → values="Value")
    gdp_per_capita.rename(columns={"2015": "GDP per capita"}, inplace=True)
    gdp_per_capita.set_index("Country", inplace=True)
    full_country_stats = pd.merge(left=oecd_bli, right=gdp_per_capita,
    left_index=True, right_index=True)
    full_country_stats.sort_values(by="GDP per capita", inplace=True)
    remove_indices = [0, 1, 6, 8, 33, 34, 35]
    keep_indices = list(set(range(36)) - set(remove_indices))
    return full_country_stats[["GDP per capita", 'Life satisfaction']].
    → iloc[keep_indices]
```

**Step 1: Load data into dataframe**

```
[4]: # Load the data
oecd_bli = pd.read_csv("data/oecd_bli_2015.csv", thousands=',')
gdp_per_capita = pd.read_csv("data/gdp_per_capita.
    ↳csv", thousands=',', delimiter='\t',
                                encoding='latin1', na_values="n/a")


# Prepare the data
country_stats = prepare_country_stats(oecd_bli, gdp_per_capita)
X = np.c_[country_stats["GDP per capita"]]
y = np.c_[country_stats["Life satisfaction"]]

# Visualize the data
country_stats.plot(kind='scatter', x="GDP per capita", y='Life satisfaction')
plt.show()

# Select a linear model
model = sklearn.linear_model.LinearRegression()

# Train the model
model.fit(X, y)

# Make a prediction for Cyprus
X_new = [[22587]] # Cyprus' GDP per capita
print(model.predict(X_new)) # outputs [[ 5.96242338]]
```



```
[[6.25984414]]
```

## Step 2: check the dimension of the table

```
[5]: print("The dimension of the table is: ", oecd_bli.shape)
```

The dimension of the table is: (2368, 17)

## Step 3: Look at the data

```
[6]: print(oecd_bli.head(5))
```

	LOCATION	Country	INDICATOR	Indicator	MEASURE	\
0	AUS	Australia	JE_LMIS	Labour market insecurity	L	
1	AUT	Austria	JE_LMIS	Labour market insecurity	L	
2	BEL	Belgium	JE_LMIS	Labour market insecurity	L	
3	CAN	Canada	JE_LMIS	Labour market insecurity	L	
4	CZE	Czech Republic	JE_LMIS	Labour market insecurity	L	

	Measure	INEQUALITY	Inequality	Unit	Code	Unit	PowerCode	Code	\
0	Value	TOT	Total	PC	Percentage		0		
1	Value	TOT	Total	PC	Percentage		0		
2	Value	TOT	Total	PC	Percentage		0		
3	Value	TOT	Total	PC	Percentage		0		
4	Value	TOT	Total	PC	Percentage		0		

	PowerCode	Reference	Period	Code	Reference	Period	Value	Flag	Codes	Flags
0	Units			NaN		NaN	5.4		NaN	NaN
1	Units			NaN		NaN	3.5		NaN	NaN
2	Units			NaN		NaN	3.7		NaN	NaN
3	Units			NaN		NaN	6.0		NaN	NaN
4	Units			NaN		NaN	3.1		NaN	NaN

```
[7]: oecd_bli.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2368 entries, 0 to 2367
Data columns (total 17 columns):
LOCATION                2368 non-null object
Country               2368 non-null object
INDICATOR              2368 non-null object
Indicator              2368 non-null object
MEASURE                2368 non-null object
Measure               2368 non-null object
INEQUALITY             2368 non-null object
Inequality             2368 non-null object
Unit Code              2368 non-null object
Unit                  2368 non-null object
```

```

PowerCode Code      2368 non-null int64
PowerCode           2368 non-null object
Reference Period Code 0 non-null float64
Reference Period     0 non-null float64
Value               2368 non-null float64
Flag Codes          0 non-null float64
Flags               0 non-null float64
dtypes: float64(5), int64(1), object(11)
memory usage: 314.6+ KB

```

Looking at the results of the “info” method, there are a number of empty columns that can be removed. There are also some with the same value throughout.

```

[8]: # Remove empty columns
oeed_bli.dropna(axis=1, inplace=True)

# Looks like some other variables have the same value from top to bottom,
# so really don't need them...
print(oeed_bli['MEASURE'].value_counts())
print(oeed_bli['PowerCode Code'].value_counts())
oeed_bli.drop(['MEASURE', 'Measure', 'PowerCode Code'], axis = 1, inplace=True)

# Remove space from Unit Code
oeed_bli.rename(columns={'Unit Code': 'UnitCode'})

print(oeed_bli.info())

```

```

L      2368
Name: MEASURE, dtype: int64
0      2368
Name: PowerCode Code, dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2368 entries, 0 to 2367
Data columns (total 10 columns):
LOCATION      2368 non-null object
Country      2368 non-null object
INDICATOR    2368 non-null object
Indicator    2368 non-null object
INEQUALITY   2368 non-null object
Inequality   2368 non-null object
Unit Code    2368 non-null object
Unit         2368 non-null object
PowerCode    2368 non-null object
Value        2368 non-null float64
dtypes: float64(1), object(9)
memory usage: 185.1+ KB
None

```

```
[9]: oecd_bli.sample(10)
```

```
[9]:      LOCATION      Country INDICATOR \
1948      GBR  United Kingdom  SW_LIFS
1916      CHE    Switzerland  SW_LIFS
1591      LVA         Latvia   HS_LEB
1061      CZE  Czech Republic  ES_STCS
2086      SVN      Slovenia   PS_REPH
547       CHE    Switzerland  JE_EMPL
1316      CHL         Chile   ES_EDUEX
192       BRA         Brazil   CG_SENG
1795      LUX      Luxembourg  HS_SFRH
1939      ISL         Iceland  SW_LIFS
```

```

Indicator INEQUALITY Inequality \
1948      Life satisfaction      HGH      High
1916      Life satisfaction      WMN      Women
1591      Life expectancy       MN       Men
1061      Student skills        MN       Men
2086      Homicide rate        WMN      Women
547       Employment rate      MN       Men
1316      Years in education    WMN      Women
192      Stakeholder engagement for developing regulations  TOT      Total
1795      Self-reported health  LW       Low
1939      Life satisfaction      HGH      High
```

```

Unit Code      Unit PowerCode  Value
1948  AVSCORE  Average score    Units    7.1
1916  AVSCORE  Average score    Units    7.5
1591      YR      Years        Units   69.8
1061  AVSCORE  Average score    Units  489.0
2086  RATIO     Ratio          Units    0.7
547    PC      Percentage      Units   84.0
1316  YR      Years          Units   17.7
192   AVSCORE  Average score    Units    2.2
1795  PC      Percentage      Units   65.0
1939  AVSCORE  Average score    Units    8.0
```

Using `pandas.pivot_table`, transform dataframe into a more human-friendly format...

```
[10]: df_table = pd.pivot_table(oecd_bli, values='Value', index='Country',
    ↪columns=['INDICATOR'])

# I also need a 'Country' column. I know this is probably not the right way to
    ↪go about this...
df_table['country'] = df_table.index.astype('str')
```

```

# Drop this row, it's not a country...
indexNames = df_table[df_table['country'] == 'OECD - Total'].index
df_table.drop(indexNames, inplace=True)

print(df_table.sample(5))

# For reference, create a dictionary of Indicators
print("LIST OF INDICATOR KEYS AND DESCRIPTIONS")
print("=====")
df_indicators = oecd_bli.groupby('INDICATOR')['Indicator'].agg('min')
print(df_indicators.sort_values())

```

INDICATOR	CG_SENG	CG_VOTO	EQ_AIRP	EQ_WATER	ES_EDUA	ES_EDUEX	\
Country							
Australia	2.7	91.0	5.0	92.666667	81.000000	20.966667	
New Zealand	2.5	80.0	5.0	89.000000	78.666667	17.700000	
Slovak Republic	3.0	60.0	21.0	84.666667	91.333333	15.766667	
Spain	1.8	70.0	11.0	72.333333	59.000000	17.900000	
Turkey	1.5	86.0	20.0	65.000000	39.000000	18.300000	

INDICATOR	ES_STCS	HO_BASE	HO_HISH	HO_NUMR	HS_LEB	HS_SFRH	\
Country							
Australia	411.2	NaN	20.0	NaN	82.500000	87.25	
New Zealand	506.2	NaN	26.0	2.4	81.700000	89.25	
Slovak Republic	463.4	1.2	23.0	1.1	77.266667	68.60	
Spain	492.4	0.1	21.0	1.9	83.400000	74.00	
Turkey	426.8	8.0	20.0	1.0	78.000000	70.40	

INDICATOR	IW_HADI	IW_HNFW	JE_EMPL	JE_LMIS	JE_LTUR	JE_PEARL	\
Country							
Australia	32759.0	427064.0	73.000000	5.922	1.306667	49126.0	
New Zealand	NaN	388514.0	77.000000	4.700	0.736667	40043.0	
Slovak Republic	20474.0	NaN	66.000000	21.376	4.773333	24328.0	
Spain	23999.0	373548.0	62.333333	23.792	7.710000	38507.0	
Turkey	NaN	NaN	51.666667	12.060	2.660000	NaN	

INDICATOR	PS_FSAFEN	PS_REPH	SC_SNTWS	SW_LIFS	WL_EWLH	WL_TNOW	\
Country							
Australia	64.133333	1.100000	95.25	7.350	12.840000	14.350000	
New Zealand	66.266667	1.300000	96.25	7.300	15.036667	14.883333	
Slovak Republic	63.700000	0.800000	91.50	6.425	4.073333	NaN	
Spain	82.166667	0.600000	92.75	6.225	3.963333	15.860000	
Turkey	59.833333	1.366667	86.00	5.520	31.043333	14.653333	

INDICATOR	country
Country	
Australia	Australia

New Zealand                      New Zealand  
 Slovak Republic    Slovak Republic  
 Spain                                      Spain  
 Turkey                                      Turkey  
 LIST OF INDICATOR KEYS AND DESCRIPTIONS

=====

INDICATOR  
 EQ\_AIRP                                      Air pollution  
 HO\_BASE                                      Dwellings without basic facilities  
 ES\_EDUA                                      Educational attainment  
 WL\_EWLH                                      Employees working very long hours  
 JE\_EMPL                                      Employment rate  
 PS\_FSAFEN                                      Feeling safe walking alone at night  
 PS\_REPH                                      Homicide rate  
 IW\_HADI                                      Household net adjusted disposable income  
 IW\_HNFW                                      Household net wealth  
 HO\_HISH                                      Housing expenditure  
 JE\_LMIS                                      Labour market insecurity  
 HS\_LEB                                      Life expectancy  
 SW\_LIFS                                      Life satisfaction  
 JE\_LTUR                                      Long-term unemployment rate  
 JE\_PEARN                                      Personal earnings  
 SC\_SNTWS                                      Quality of support network  
 HO\_NUMR                                      Rooms per person  
 HS\_SFRH                                      Self-reported health  
 CG\_SENG                                      Stakeholder engagement for developing regulations  
 ES\_STCS                                      Student skills  
 WL\_TNOW                                      Time devoted to leisure and personal care  
 CG\_VOTO                                      Voter turnout  
 EQ\_WATER                                      Water quality  
 ES\_EDUEX                                      Years in education  
 Name: Indicator, dtype: object

```
[11]: print("Describe Data")
      print(df_table.describe())
```

```
Describe Data
INDICATOR    CG_SENG    CG_VOTO    EQ_AIRP    EQ_WATER    ES_EDUA    ES_EDUEX  \
count      38.000000   40.00000   40.000000  40.000000   39.000000  39.000000
mean        2.160526   69.57500   13.325000  82.333333   77.717949  17.547863
std         0.577291   12.21157    5.770782  10.492977   15.136134    1.412720
min         1.200000   47.00000    3.000000  55.333333   37.666667  14.100000
25%         1.725000   60.75000    9.750000  74.250000   75.000000  16.550000
50%         2.200000   69.50000   14.000000  83.833333   82.000000  17.666667
75%         2.575000   79.00000   16.500000  91.083333   87.833333  18.350000
max         3.200000   91.00000   28.000000  98.666667   94.000000  20.966667
```

INDICATOR	ES_STCS	HO_BASE	HO_HISH	HO_NUMR	HS_LEB	HS_SFRH \
count	39.000000	37.000000	38.000000	37.000000	40.000000	37.000000
mean	485.707692	5.075676	20.657895	1.632432	79.567500	67.493243
std	33.787972	8.448320	2.528500	0.431441	4.669642	14.331584
min	398.200000	0.000000	15.000000	0.900000	57.500000	33.000000
25%	475.800000	0.300000	19.000000	1.200000	77.916667	60.800000
50%	492.800000	0.900000	21.000000	1.600000	81.366667	70.200000
75%	506.800000	6.700000	22.750000	1.900000	82.366667	76.000000
max	528.800000	37.000000	26.000000	2.600000	84.066667	89.250000

INDICATOR	IW_HADI	IW_HNFW	JE_EMPL	JE_LMIS	JE_LTUR \
count	29.000000	27.000000	40.000000	33.000000	38.000000
mean	27807.310345	289780.185185	68.533333	7.706970	2.855789
std	7055.262661	165673.432787	7.882253	6.234572	3.622899
min	16275.000000	70160.000000	43.333333	0.662000	0.050000
25%	21453.000000	180100.000000	65.833333	4.392000	1.011667
50%	29333.000000	259667.000000	69.666667	5.396000	1.776667
75%	31304.000000	379777.000000	74.000000	8.784000	3.196667
max	45284.000000	769053.000000	85.666667	29.200000	16.643333

INDICATOR	JE_PEARL	PS_FSAFEN	PS_REPH	SC_SNTWS	SW_LIFS \
count	35.000000	40.000000	40.000000	40.000000	40.000000
mean	39817.514286	68.463333	3.481667	90.193333	6.577208
std	13108.329748	13.960934	6.459861	4.384954	0.762724
min	15314.000000	35.866667	0.166667	78.333333	4.700000
25%	25971.500000	60.108333	0.600000	88.300000	5.938333
50%	40863.000000	70.483333	0.950000	91.350000	6.510000
75%	49400.500000	78.500000	2.166667	93.062500	7.243750
max	63062.000000	90.033333	27.000000	98.000000	7.660000

INDICATOR	WL_EWLH	WL_TNOW
count	38.000000	22.000000
mean	7.789649	15.048939
std	7.585983	0.672978
min	0.140000	13.826667
25%	3.150833	14.560833
50%	4.981667	14.885000
75%	10.571667	15.600833
max	31.043333	16.336667

```
[12]: corr_matrix = df_table.corr()
corr_matrix["SW_LIFS"].sort_values(ascending=False)
```

```
[12]: INDICATOR
SW_LIFS      1.000000
JE_PEARL     0.731418
IW_HADI      0.713008
EQ_WATER     0.682587
```



JE_EMPL	0.678344
SC_SNTWS	0.667896
HS_SFRH	0.656817
PS_FSAFEN	0.600163
HO_NUMR	0.597502
HS_LEB	0.568044
CG_VOTO	0.368598
ES_EDUEX	0.324655
ES_EDUA	0.293395
IW_HNFW	0.292887
HO_HISH	0.286334
WL_TNOW	0.199424
ES_STCS	0.197223
CG_SENG	0.180861
WL_EWLH	-0.195136
PS_REPH	-0.259378
JE_LMIS	-0.452874
HO_BASE	-0.528167
EQ_AIRP	-0.551376
JE_LTUR	-0.567002

Name: SW\_LIFS, dtype: float64

**Step 4: Think about some questions that might help you predict what indicators most influence the Life Satisfaction score:** The central point of this dataset is the so-called, “**Life Satisfaction Index**”. In other words, do indicators in the categories of housing, income, jobs, community, education, environment, civic engagement, health, etc. really lead to a better, more satisfied life? Let’s focus on a few high-level categories to see how the indicators correlate with the LSI...

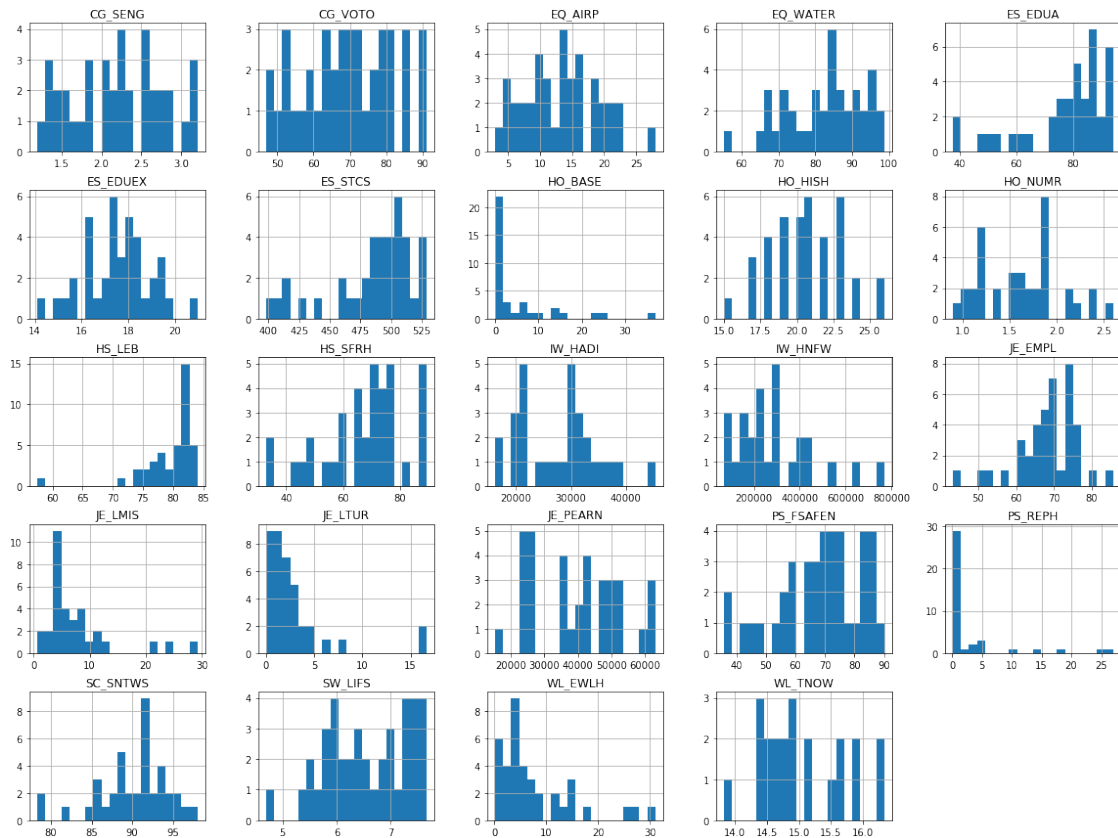
- **Wealth** Net Wealth, Labor Market Insecurity, Employment rate
- **Environment** Air pollution, Homicide rate, Water quality
- **Health** Life expectancy, Self-reported health, Long work hours

My observations are based on the **graph analysis** below.

First, there does not seem to be any noticeable normal distribution amongst any of the indicators, although some—such as HS\_LEB (Life Expectancy) exhibit *normal-ish* distribution on a skewed scale.

- **Wealth** - somewhat surprisingly, Net Wealth does not appear to be as important as labor market security and the employment rate. Having said that, removing the data points above \$500,000 might tell a different story.
- **Environment** - Air and water quality seem to factor higher than the homicide rate, which shows almost no effect on the LSI.
- **Health** - Life expectancy seems like an obvious factor, but I was also satisfied to see long work hours affect the index as well.

```
[13]: df_table.hist( bins = 20, figsize =( 20,15))
plt.show()
```



```
[14]: results = smf.ols('SW_LIFS ~ IW_HNFW + JE_LMIS + JE_EMPL', data=df_table).fit()
print("Money: Net Wealth, Labor Market Insecurity, Employment rate")
print(results.summary())
```

Money: Net Wealth, Labor Market Insecurity, Employment rate

#### OLS Regression Results

```
=====
Dep. Variable:          SW_LIFS    R-squared:                0.327
Model:                  OLS        Adj. R-squared:           0.235
Method:                 Least Squares    F-statistic:            3.564
Date:                  Sun, 26 Jan 2020    Prob (F-statistic):      0.0306
Time:                  09:45:22    Log-Likelihood:         -22.710
No. Observations:      26    AIC:                    53.42
Df Residuals:          22    BIC:                    58.45
Df Model:              3
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.7241	2.217	1.229	0.232	-1.873	7.321
IW_HNFW	1.13e-06	7.6e-07	1.486	0.151	-4.47e-07	2.71e-06

JE_LMIS	-0.0085	0.028	-0.304	0.764	-0.067	0.050
JE_EMPL	0.0523	0.030	1.763	0.092	-0.009	0.114
=====						
Omnibus:		0.953	Durbin-Watson:			1.785
Prob(Omnibus):		0.621	Jarque-Bera (JB):			0.831
Skew:		-0.167	Prob(JB):			0.660
Kurtosis:		2.190	Cond. No.			6.02e+06
=====						


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The condition number is large, 6.02e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
[15]: attributes = ['SW_LIFS', 'IW_HNFW', 'JE_LMIS', 'JE_EMPL']
scatter_matrix(df_table[attributes], alpha=1.0, figsize=(12, 8))
```

```
[15]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6B1179A0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6B0E66D0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A9D3B50>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A9FB3D0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AA1CC10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AA4B490>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AA563A0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AA81C40>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AAD2D90>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AB07610>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AB32E50>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AB666D0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AB91F10>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6ABC5790>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6ABF1FD0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AC26850>]],
dtype=object)
```



```
[16]: (ggplot(df_table, aes(x='IW_HNFW', y='SW_LIFS', color='JE_LMIS',
→size='JE_EMPL')) +
      geom_point(alpha=0.5) +
      scale_x_log10() +
      geom_text(aes(x='IW_HNFW', y='SW_LIFS', label='country'),
                color="black",
                size=8,
                data=df_table) +
      theme(figure_size = (12.0, 8.0)) +
      labs(title="Life Satisfaction versus Wealth Factors",x="Net_
→Wealth",y="Life Satisfaction Index",size="Employment Rate",color="Labor_
→Market Insecurity")
    )
```




[16]: <ggplot: (174595026035)>

```
[17]: import yellowbrick
from yellowbrick.features import Rank2D
from yellowbrick.features import ParallelCoordinates
from yellowbrick.style import set_palette

#set up the figure size
plt.rcParams['figure.figsize'] = (15, 7)
num_features = ['SW_LIFS', 'IW_HNFW', 'JE_LMIS', 'JE_EMPL']
# extract the numpy arrays from the data frame
X = df_table[num_features].as_matrix()

# instantiate the visualizer with the Covariance ranking algorithm
visualizer = Rank2D(features=num_features, algorithm='pearson')
visualizer.fit(X) # Fit the data to the visualizer
visualizer.transform(X) # Transform the data

plt.show()
```



```
[18]: results = smf.ols('SW_LIFS ~ EQ_AIRP + PS_REPH + EQ_WATER', data=df_table).fit()
print("Environment: Air pollution, Homicide rate, Water quality")
print(results.summary())
```

Environment: Air pollution, Homicide rate, Water quality

#### OLS Regression Results

```
=====
Dep. Variable:          SW_LIFS    R-squared:                0.512
Model:                  OLS        Adj. R-squared:           0.471
Method:                 Least Squares    F-statistic:           12.57
Date:                   Sun, 26 Jan 2020    Prob (F-statistic):    8.92e-06
Time:                   09:45:24    Log-Likelihood:       -31.080
No. Observations:       40    AIC:                   70.16
Df Residuals:           36    BIC:                   76.92
Df Model:                3
Covariance Type:        nonrobust
=====
```

```
=====
coef    std err          t    P>|t|    [0.025    0.975]
-----
```

Intercept	3.8394	1.218	3.152	0.003	1.369	6.310
EQ_AIRP	-0.0343	0.020	-1.751	0.088	-0.074	0.005
PS_REPH	-0.0015	0.017	-0.093	0.926	-0.035	0.032
EQ_WATER	0.0389	0.012	3.193	0.003	0.014	0.064

```
=====
Omnibus:                3.892    Durbin-Watson:                1.767
Prob(Omnibus):          0.143    Jarque-Bera (JB):        2.842
Skew:                   -0.368    Prob(JB):                0.241
Kurtosis:               4.079    Cond. No.                1.17e+03
=====
```


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.17e+03. This might indicate that there are strong multicollinearity or other numerical problems.


```
[19]: attributes = ['SW_LIFS', 'EQ_AIRP', 'PS_REPH', 'EQ_WATER']
      scatter_matrix(df_table[attributes], alpha=1.0, figsize=(12, 8))
```

```
[19]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A85F520>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A6A38B0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A73BE20>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A797B20>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A6D05B0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6B00C6D0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6B00CDC0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A71DB50>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A6394C0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A3EA250>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A6570D0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A6EF550>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A247C70>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6B01B8B0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A8FD400>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A934670>]],
      dtype=object)
```



```
[20]: (ggplot(df_table, aes(x='EQ_AIRP', y='SW_LIFS', color='PS_REPH',
→size='EQ_WATER')) +
  geom_point(alpha = 0.5) +
  geom_text(aes(x='EQ_AIRP', y='SW_LIFS', label='country'),
    color="black",
    size=8,
    data=df_table) +
  theme(figure_size = (12.0, 8.0)) +
  labs(title="Life Satisfaction versus Environment",x="Air_
→Quality",y="Life Satisfaction Index",size="Water Quality",color="Homicide_
→Rate")
)
```






[20]: <ggplot: (174595291127)>

```
[21]: #set up the figure size
plt.rcParams['figure.figsize'] = (15, 7)
num_features = ['SW_LIFS', 'EQ_AIRP', 'PS_REPH', 'EQ_WATER']
# extract the numpy arrays from the data frame
X = df_table[num_features].as_matrix()

# instantiate the visualizer with the Covariance ranking algorithm
visualizer = Rank2D(features=num_features, algorithm='pearson')
visualizer.fit(X) # Fit the data to the visualizer
visualizer.transform(X) # Transform the data

plt.show()
```



```
[22]: results = smf.ols('SW_LIFS ~ HS_LEB + HS_SFRH + WL_EWLH', data=df_table).fit()
print("Health: Life expectancy, Self-reported health, Long work hours")
print(results.summary())
```

Health: Life expectancy, Self-reported health, Long work hours

#### OLS Regression Results

```
=====
Dep. Variable:          SW_LIFS    R-squared:                0.484
Model:                  OLS        Adj. R-squared:           0.434
Method:                 Least Squares    F-statistic:             9.707
Date:                  Sun, 26 Jan 2020    Prob (F-statistic):       0.000114
Time:                  09:45:26          Log-Likelihood:          -26.111
No. Observations:      35              AIC:                    60.22
Df Residuals:          31              BIC:                    66.44
Df Model:              3
Covariance Type:       nonrobust
=====
```

```
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
```

Intercept	2.3807	3.176	0.750	0.459	-4.096	8.858
HS_LEB	0.0227	0.046	0.499	0.621	-0.070	0.116
HS_SFRH	0.0382	0.011	3.403	0.002	0.015	0.061
WL_EWLH	-0.0250	0.014	-1.737	0.092	-0.054	0.004


```
=====
Omnibus:                8.285    Durbin-Watson:                2.096
Prob(Omnibus):          0.016    Jarque-Bera (JB):        6.935
Skew:                   -0.902    Prob(JB):                0.0312
Kurtosis:               4.227    Cond. No.                3.70e+03
=====
```

Warnings:


- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 3.7e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
[23]: attributes = ['SW_LIFS', 'HS_LEB', 'HS_SFRH', 'WL_EWLH']
      scatter_matrix(df_table[attributes], alpha=1.0, figsize=(12, 8))
```

```
[23]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A70AA30>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A429EE0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A68D3D0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A4C9850>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A895CD0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A509040>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A5090A0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A5CA4C0>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A53AF10>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A4383D0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A434850>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A7C2400>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A8B5D00>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AD24FD0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6A3CB790>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x0000028A6AE55580>]],
      dtype=object)
```



```
[24]: (ggplot(df_table[df_table['HS_LEB'] > 70], aes(x='HS_LEB', y='SW_LIFS',
→color='WL_EWLH', size='HS_SFRH')) +
  geom_point(alpha = 0.8) +
  scale_x_continuous(limits=[70,85]) +
  geom_text(aes(x='HS_LEB', y='SW_LIFS', label='country'),
    color="black",
    size=8,
    data=df_table) +
  theme(figure_size = (12.0, 8.0)) +
  labs(title="Life Satisfaction versus Health Factors",x="Life_
→Expectancy",y="Life Satisfaction Index",size="Self-reported_
→Health",color="Long Work Hours")
)
```




[24]: <ggplot: (174594488278)>

```
[25]: #set up the figure size
plt.rcParams['figure.figsize'] = (15, 7)
num_features = ['SW_LIFS', 'HS_LEB', 'HS_SFRH', 'WL_EWLH']
# extract the numpy arrays from the data frame
X = df_table[num_features].as_matrix()

# instantiate the visualizer with the Covariance ranking algorithm
visualizer = Rank2D(features=num_features, algorithm='pearson')
visualizer.fit(X) # Fit the data to the visualizer
visualizer.transform(X) # Transform the data

plt.show()
```



## Part 2: Dimensionality and Feature Reduction

[26]: *# What kind of data are we dealing with?*  
df\_table.describe()

[26]:

INDICATOR	CG_SENG	CG_VOTO	EQ_AIRP	EQ_WATER	ES_EDUA	ES_EDUEX	\
count	38.000000	40.00000	40.000000	40.000000	39.000000	39.000000	
mean	2.160526	69.57500	13.325000	82.333333	77.717949	17.547863	
std	0.577291	12.21157	5.770782	10.492977	15.136134	1.412720	
min	1.200000	47.00000	3.000000	55.333333	37.666667	14.100000	
25%	1.725000	60.75000	9.750000	74.250000	75.000000	16.550000	
50%	2.200000	69.50000	14.000000	83.833333	82.000000	17.666667	
75%	2.575000	79.00000	16.500000	91.083333	87.833333	18.350000	
max	3.200000	91.00000	28.000000	98.666667	94.000000	20.966667	

INDICATOR	ES_STCS	HO_BASE	HO_HISH	HO_NUMR	HS_LEB	HS_SFRH	\
count	39.000000	37.000000	38.000000	37.000000	40.000000	37.000000	
mean	485.707692	5.075676	20.657895	1.632432	79.567500	67.493243	
std	33.787972	8.448320	2.528500	0.431441	4.669642	14.331584	
min	398.200000	0.000000	15.000000	0.900000	57.500000	33.000000	
25%	475.800000	0.300000	19.000000	1.200000	77.916667	60.800000	
50%	492.800000	0.900000	21.000000	1.600000	81.366667	70.200000	
75%	506.800000	6.700000	22.750000	1.900000	82.366667	76.000000	
max	528.800000	37.000000	26.000000	2.600000	84.066667	89.250000	

INDICATOR	IW_HADI	IW_HNFW	JE_EMPL	JE_LMIS	JE_LTUR	\
count	29.000000	27.000000	40.000000	33.000000	38.000000	
mean	27807.310345	289780.185185	68.533333	7.706970	2.855789	
std	7055.262661	165673.432787	7.882253	6.234572	3.622899	
min	16275.000000	70160.000000	43.333333	0.662000	0.050000	
25%	21453.000000	180100.000000	65.833333	4.392000	1.011667	
50%	29333.000000	259667.000000	69.666667	5.396000	1.776667	
75%	31304.000000	379777.000000	74.000000	8.784000	3.196667	
max	45284.000000	769053.000000	85.666667	29.200000	16.643333	

INDICATOR	JE_PEARL	PS_FSAFEN	PS_REPH	SC_SNTWS	SW_LIFS	\
count	35.000000	40.000000	40.000000	40.000000	40.000000	
mean	39817.514286	68.463333	3.481667	90.193333	6.577208	
std	13108.329748	13.960934	6.459861	4.384954	0.762724	
min	15314.000000	35.866667	0.166667	78.333333	4.700000	
25%	25971.500000	60.108333	0.600000	88.300000	5.938333	
50%	40863.000000	70.483333	0.950000	91.350000	6.510000	
75%	49400.500000	78.500000	2.166667	93.062500	7.243750	
max	63062.000000	90.033333	27.000000	98.000000	7.660000	

INDICATOR	WL_EWLH	WL_TNOW
count	38.000000	22.000000
mean	7.789649	15.048939
std	7.585983	0.672978
min	0.140000	13.826667
25%	3.150833	14.560833
50%	4.981667	14.885000
75%	10.571667	15.600833
max	31.043333	16.336667

### Proposed Steps for Dimensionality and Feature Reduction

- Due to the quantitative nature of these variables, there exists wide discrepancies of units. For instance, Household Net Wealth can be in the thousands of US dollars, while Employment Rate is a percentage. I will apply a simple min-max rescaler to normalize all variables.
- The *variance* between features varies greatly, For example, CG\_SENS ranges from 1.2 to 3.2 while SC\_SNTWS ranges from 4.38 to 98. I will start by setting a variance threshold and eliminating features below the threshold. In theory, this should eliminate variables with low

variance, which likely will not contribute greatly to the model.

- I suspect that variables such as air pollution and water quality probably are highly correlated. I will apply a correlation matrix and will consider dropping one of the correlated features.
- That takes care of the easy stuff. Next I'd like to automatically select the best features to keep by leveraging scikit-learn's recursive feature elimination functionality.

#### VarianceThreshold

```
[27]: from sklearn import preprocessing

# Make a features dataset by dropping the target variable--SW_LIFS
features = df_table.drop(['country', 'SW_LIFS'], axis=1)
target = df_table['SW_LIFS']

# Standardizing
std_scale = preprocessing.StandardScaler().fit(features)
df_std = std_scale.transform(features)

#Min-max scaling
minmax_scale = preprocessing.MinMaxScaler().fit(features)
df_minmax = minmax_scale.transform(features)

[28]: def plot():
    plt.figure(figsize=(8,6))

    plt.scatter(df_table['JE_EMPL'], df_table['IW_HNFW'],
                color='green', label='original scale', alpha=0.5)

    plt.scatter(df_std[:,4], df_std[:,13], color='red',
                label='standardized [mu=0,sigma=1]', alpha=0.3)


    plt.scatter(df_minmax[:,4], df_minmax[:,13],
                color='blue', label='min-max scaled [min=0, max=1]', alpha=0.3)

    plt.title('Employment Rate and Household Income')
    plt.xlabel('Empl Rate')
    plt.ylabel('HH Income')
    plt.legend(loc='upper left')
    plt.xscale('symlog')
    plt.yscale('symlog')
    plt.grid()

    plt.tight_layout()

plot()
plt.show()
```





```
[29]: from sklearn.feature_selection import VarianceThreshold

# Make a features dataset by dropping the target variable--SW_LIFS
features = df_table.drop(['country', 'SW_LIFS'], axis=1)
target = df_table['SW_LIFS']

# Create thresholder
thresholder = VarianceThreshold(threshold=5.0)

# Create high variance feature matrix
features_high_variance = thresholder.fit_transform(features)

# View high variance feature matrix
features_high_variance[0:3]

# View variances
thresholder.fit(features).variances_
```

```
[29]: array([3.24494460e-01, 1.45394375e+02, 3.24693750e+01, 1.07350000e+02,
        2.23228139e+02, 1.94460370e+00, 1.11235456e+03, 6.94450840e+01,
        6.22506925e+00, 1.81110299e-01, 2.12604160e+01, 1.99843100e+02,
        4.80602922e+07, 2.64311054e+10, 6.05766667e+01, 3.76920085e+01,
```

```
1.27799922e+01, 1.66918929e+08, 1.90034989e+02, 4.06865528e+01,
1.87471222e+01, 5.60327426e+01, 4.32312511e-01])
```

```
[30]: # Create correlation matrix
corr_matrix = features.corr().abs()

# Select upper triangle of correlation matrix
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape),
                                     k=1).astype(np.bool))

# Find index of feature columns with correlation greater than 0.95
to_drop = [column for column in upper.columns if any(upper[column] > 0.95)]

# Drop features
features.drop(features.columns[to_drop], axis=1).head()
```

```
[30]: INDICATOR  CG_SENG  CG_VOTO  EQ_AIRP   EQ_WATER   ES_EDUA   ES_EDUEX  \
Country
Australia      2.7      91.0      5.0  92.666667  81.000000  20.966667
Austria         1.3      80.0     16.0  92.000000  85.000000  17.000000
Belgium         2.0      89.0     15.0  83.666667  77.000000  19.300000
Brazil          2.2      79.0     10.0  73.000000  49.000000  16.166667
Canada          2.9      68.0      7.0  91.000000  91.333333  17.333333
```

```
INDICATOR  ES_STCS  HO_BASE  HO_HISH  HO_NUMR   HS_LEB  HS_SFRH  IW_HADI  \
Country
Australia  411.2    NaN      20.0     NaN  82.500000  87.25  32759.0
Austria    492.8    0.9     21.0     1.6  81.700000  70.60  33541.0
Belgium    503.8    1.9     21.0     2.2  81.500000  73.60  30364.0
Brazil     398.2    6.7     NaN     NaN  74.766667  NaN    NaN
Canada     523.2    0.2     22.0     2.6  81.866667  87.80  30854.0
```

```
INDICATOR  IW_HNFW   JE_EMPL  JE_LMIS   JE_LTUR  JE_PEARL  PS_FSAFEN  \
Country
Australia  427064.0  73.000000  5.922  1.306667  49126.0  64.133333
Austria    308325.0  72.000000  4.076  1.830000  50349.0  80.700000
Belgium    386006.0  63.333333  4.052  3.533333  49675.0  70.266667
Brazil      NaN  61.000000  NaN     NaN     NaN     NaN  35.866667
Canada     423849.0  73.333333  7.048  0.763333  47622.0  82.500000
```

```
INDICATOR  PS_REPH  SC_SNTWS   WL_EWLH   WL_TNOW
Country
Australia  1.100000  95.25  12.840000  14.350000
Austria    0.466667  92.00  6.590000  14.530000
Belgium    1.033333  92.00  4.703333  15.663333
Brazil     27.000000  89.25  7.006667  NaN
Canada     1.266667  93.25  3.673333  14.553333
```

```
[31]: import seaborn as sns
print(corr_matrix)
sns.heatmap(corr_matrix)
```

INDICATOR	CG_SENG	CG_VOTO	EQ_AIRP	EQ_WATER	ES_EDUA	ES_EDUEX	\
INDICATOR							
CG_SENG	1.000000	0.028002	0.015333	0.052708	0.245220	0.037028	
CG_VOTO	0.028002	1.000000	0.205751	0.161888	0.086394	0.301847	
EQ_AIRP	0.015333	0.205751	1.000000	0.545403	0.012111	0.341326	
EQ_WATER	0.052708	0.161888	0.545403	1.000000	0.401520	0.499011	
ES_EDUA	0.245220	0.086394	0.012111	0.401520	1.000000	0.272450	
ES_EDUEX	0.037028	0.301847	0.341326	0.499011	0.272450	1.000000	
ES_STCS	0.067136	0.057214	0.072012	0.473621	0.627987	0.248241	
HO_BASE	0.002465	0.231873	0.229820	0.586296	0.327014	0.500586	
HO_HISH	0.070164	0.046796	0.204556	0.259928	0.052753	0.146336	
HO_NUMR	0.132029	0.297270	0.599746	0.632610	0.238163	0.323064	
HS_LEB	0.086465	0.147134	0.287432	0.502301	0.086281	0.314825	
HS_SFRH	0.186154	0.422847	0.416346	0.295603	0.028642	0.257713	
IW_HADI	0.009595	0.450701	0.384807	0.474136	0.069743	0.062532	
IW_HNFW	0.313337	0.399376	0.197628	0.076700	0.030136	0.304318	
JE_EMPL	0.149286	0.001314	0.525146	0.726957	0.504685	0.238336	
JE_LMIS	0.046940	0.272671	0.203927	0.456311	0.236128	0.007565	
JE_LTUR	0.199630	0.063474	0.284824	0.390481	0.202390	0.132076	
JE_PEARN	0.019460	0.529236	0.450445	0.604712	0.218113	0.245365	
PS_FSAFEN	0.042428	0.152899	0.364504	0.745049	0.449514	0.456138	
PS_REPH	0.041994	0.136966	0.007333	0.464015	0.490610	0.503941	
SC_SNTWS	0.043415	0.297239	0.679765	0.622848	0.276437	0.299917	
WL_EWLH	0.009313	0.067688	0.115445	0.376035	0.634258	0.282485	
WL_TNOW	0.302802	0.307280	0.088929	0.004498	0.337958	0.094854	

INDICATOR	ES_STCS	HO_BASE	HO_HISH	HO_NUMR	HS_LEB	HS_SFRH	\
INDICATOR							
CG_SENG	0.067136	0.002465	0.070164	0.132029	0.086465	0.186154	
CG_VOTO	0.057214	0.231873	0.046796	0.297270	0.147134	0.422847	
EQ_AIRP	0.072012	0.229820	0.204556	0.599746	0.287432	0.416346	
EQ_WATER	0.473621	0.586296	0.259928	0.632610	0.502301	0.295603	
ES_EDUA	0.627987	0.327014	0.052753	0.238163	0.086281	0.028642	
ES_EDUEX	0.248241	0.500586	0.146336	0.323064	0.314825	0.257713	
ES_STCS	1.000000	0.600704	0.126796	0.627171	0.436038	0.171635	
HO_BASE	0.600704	1.000000	0.359024	0.570377	0.844605	0.480710	
HO_HISH	0.126796	0.359024	1.000000	0.194369	0.277640	0.329240	
HO_NUMR	0.627171	0.570377	0.194369	1.000000	0.587021	0.464200	
HS_LEB	0.436038	0.844605	0.277640	0.587021	1.000000	0.404711	
HS_SFRH	0.171635	0.480710	0.329240	0.464200	0.404711	1.000000	
IW_HADI	0.028382	0.451797	0.095344	0.717301	0.459078	0.515910	
IW_HNFW	0.045354	0.428187	0.191833	0.571859	0.374210	0.355329	
JE_EMPL	0.478432	0.489324	0.215731	0.435517	0.522467	0.085938	


JE_LMIS	0.288360	0.063078	0.164635	0.264569	0.122740	0.125001
JE_LTUR	0.102944	0.292733	0.037768	0.173907	0.439576	0.050118
JE_PEARN	0.299373	0.596020	0.080009	0.711735	0.641399	0.513532
PS_FSAFEN	0.739562	0.765774	0.275512	0.655590	0.690902	0.345593
PS_REPH	0.664774	0.707629	0.364861	0.368320	0.585053	0.183617
SC_SNTWS	0.284479	0.383604	0.376605	0.606296	0.278339	0.525843
WL_EWLH	0.630973	0.534097	0.033516	0.274522	0.248687	0.252242
WL_TNOW	0.188257	0.318026	0.101389	0.193615	0.320410	0.147207

INDICATOR	IW_HADI	IW_HNFW	JE_EMPL	JE_LMIS	JE_LTUR	JE_PEARN	\
INDICATOR							
CG_SENG	0.009595	0.313337	0.149286	0.046940	0.199630	0.019460	
CG_VOTO	0.450701	0.399376	0.001314	0.272671	0.063474	0.529236	
EQ_AIRP	0.384807	0.197628	0.525146	0.203927	0.284824	0.450445	
EQ_WATER	0.474136	0.076700	0.726957	0.456311	0.390481	0.604712	
ES_EDUA	0.069743	0.030136	0.504685	0.236128	0.202390	0.218113	
ES_EDUEX	0.062532	0.304318	0.238336	0.007565	0.132076	0.245365	
ES_STCS	0.028382	0.045354	0.478432	0.288360	0.102944	0.299373	
HO_BASE	0.451797	0.428187	0.489324	0.063078	0.292733	0.596020	
HO_HISH	0.095344	0.191833	0.215731	0.164635	0.037768	0.080009	
HO_NUMR	0.717301	0.571859	0.435517	0.264569	0.173907	0.711735	
HS_LEB	0.459078	0.374210	0.522467	0.122740	0.439576	0.641399	
HS_SFRH	0.515910	0.355329	0.085938	0.125001	0.050118	0.513532	
IW_HADI	1.000000	0.674728	0.378516	0.430542	0.415803	0.917182	
IW_HNFW	0.674728	1.000000	0.099342	0.181178	0.210344	0.654348	
JE_EMPL	0.378516	0.099342	1.000000	0.614939	0.735549	0.481038	
JE_LMIS	0.430542	0.181178	0.614939	1.000000	0.857796	0.366901	
JE_LTUR	0.415803	0.210344	0.735549	0.857796	1.000000	0.258938	
JE_PEARN	0.917182	0.654348	0.481038	0.366901	0.258938	1.000000	
PS_FSAFEN	0.448296	0.211345	0.634678	0.226821	0.354233	0.694537	
PS_REPH	0.034060	0.056629	0.344470	0.037982	0.082079	0.392847	
SC_SNTWS	0.510536	0.376644	0.552591	0.286091	0.294198	0.644078	
WL_EWLH	0.187986	0.432014	0.364282	0.061697	0.004828	0.132684	
WL_TNOW	0.197591	0.089883	0.087340	0.214011	0.188335	0.237182	

INDICATOR	PS_FSAFEN	PS_REPH	SC_SNTWS	WL_EWLH	WL_TNOW
INDICATOR					
CG_SENG	0.042428	0.041994	0.043415	0.009313	0.302802
CG_VOTO	0.152899	0.136966	0.297239	0.067688	0.307280
EQ_AIRP	0.364504	0.007333	0.679765	0.115445	0.088929
EQ_WATER	0.745049	0.464015	0.622848	0.376035	0.004498
ES_EDUA	0.449514	0.490610	0.276437	0.634258	0.337958
ES_EDUEX	0.456138	0.503941	0.299917	0.282485	0.094854
ES_STCS	0.739562	0.664774	0.284479	0.630973	0.188257
HO_BASE	0.765774	0.707629	0.383604	0.534097	0.318026
HO_HISH	0.275512	0.364861	0.376605	0.033516	0.101389
HO_NUMR	0.655590	0.368320	0.606296	0.274522	0.193615
HS_LEB	0.690902	0.585053	0.278339	0.248687	0.320410

HS_SFRH	0.345593	0.183617	0.525843	0.252242	0.147207
IW_HADI	0.448296	0.034060	0.510536	0.187986	0.197591
IW_HNFW	0.211345	0.056629	0.376644	0.432014	0.089883
JE_EMPL	0.634678	0.344470	0.552591	0.364282	0.087340
JE_LMIS	0.226821	0.037982	0.286091	0.061697	0.214011
JE_LTUR	0.354233	0.082079	0.294198	0.004828	0.188335
JE_PEARN	0.694537	0.392847	0.644078	0.132684	0.237182
PS_FSAFEN	1.000000	0.739129	0.561306	0.437147	0.186304
PS_REPH	0.739129	1.000000	0.251516	0.442530	0.317128
SC_SNTWS	0.561306	0.251516	1.000000	0.246698	0.295831
WL_EWLH	0.437147	0.442530	0.246698	1.000000	0.263090
WL_TNOW	0.186304	0.317128	0.295831	0.263090	1.000000

[31]: <matplotlib.axes.\_subplots.AxesSubplot at 0x28a6b17a790>



```
[32]: # Load libraries
import warnings
from sklearn.datasets import make_regression
from sklearn.feature_selection import RFECV
from sklearn import linear_model

# Suppress an annoying but harmless warning
warnings.filterwarnings(action="ignore", module="scipy",
                        message="^internal gelsd")

# Generate features matrix, target vector, and the true coefficients
features, target = make_regression(n_samples = 10000,
```

```

n_features = 100,
n_informative = 2,
random_state = 1)

# Create a linear regression
ols = linear_model.LinearRegression()

# Recursively eliminate features
rfecv = RFECV(estimator=ols, step=1, scoring="neg_mean_squared_error")
rfecv.fit(features, target)
rfecv.transform(features)

# Once we have conducted RFE, we can see the number of features we should keep:
# Number of best features
print("Number of features we should keep: {}".format(rfecv.n_features_))

# We can also see which of those features we should keep:
# Which categories are best
print(rfecv.support_)

# Rank features best (1) to worst
print(rfecv.ranking_)

```

Number of features we should keep: 9

```

[False  True False False False  True False False False False False False
 False False False False False False  True False False False False False
 False False False False False False False False False False False True False
 False False False False False False False False False False False False
 False False False False False False False False False False False False
 False False True  True False False  True False False False False False
 False False False False]
[43  1 52 40 25  1 30 12 72  5 75 24 55 45 39 73 47 36  1 38 74 87 21 17
 13 56 11  4  3 33  9 59 22 29  1 63 34 41 84  1 10 26 28 71 78 42 91  1
 88 92 85 54 80 81 31 86 48  7 20 62 83 50  6 37 60 65 57 76 46 49 44  2
 15 66 16 35 90 82 77 69 64 32 18 51 23 67  1  1  8 53  1 89 68 58 79 61
 27 70 14 19]

```

---

## References

[http://sebastianraschka.com/Articles/2014\\_about\\_feature\\_scaling.html#standardization-and-min-max-scaling](http://sebastianraschka.com/Articles/2014_about_feature_scaling.html#standardization-and-min-max-scaling) Albon, Chris. Machine Learning with Python Cookbook: Practical Solutions from Preprocessing to Deep Learning . O'Reilly Media. Kindle Edition.