

Kinney_DSC550_Final

February 1, 2020

0.1 D. Kinney DSC 550 Final Project

0.1.1 Part 1: Graph Analysis

```
[46]: import warnings
warnings.filterwarnings("ignore")

[47]: import matplotlib.pyplot as plt
import numpy as np
import pandas as pd
import sklearn.linear_model
import statsmodels.formula.api as smf

from pandas.plotting import scatter_matrix
from plotnine import *

pd.set_option('display.max_columns', None)

%matplotlib inline

[48]: def prepare_country_stats(oecd_bli, gdp_per_capita):
    oecd_bli = oecd_bli[oecd_bli["INEQUALITY"]=="TOT"]
    oecd_bli = oecd_bli.pivot(index="Country", columns="Indicator",
    → values="Value")
    gdp_per_capita.rename(columns={"2015": "GDP per capita"}, inplace=True)
    gdp_per_capita.set_index("Country", inplace=True)
    full_country_stats = pd.merge(left=oecd_bli, right=gdp_per_capita,
    left_index=True, right_index=True)
    full_country_stats.sort_values(by="GDP per capita", inplace=True)
    remove_indices = [0, 1, 6, 8, 33, 34, 35]
    keep_indices = list(set(range(36)) - set(remove_indices))
    return full_country_stats[["GDP per capita", 'Life satisfaction']].
    → iloc[keep_indices]
```

Step 1: Load data into dataframe

```
[49]: # Load the data
oecd_bli = pd.read_csv("data/oecd_bli_2015.csv", thousands=',')
gdp_per_capita = pd.read_csv("data/gdp_per_capita.
    ↪csv", thousands=',', delimiter='\t',
                                encoding='latin1', na_values="n/a")

# Prepare the data
country_stats = prepare_country_stats(oecd_bli, gdp_per_capita)
X = np.c_[country_stats["GDP per capita"]]
y = np.c_[country_stats["Life satisfaction"]]


# Visualize the data
country_stats.plot(kind='scatter', x="GDP per capita", y='Life satisfaction')
plt.show()

# Select a linear model
model = sklearn.linear_model.LinearRegression()

# Train the model
model.fit(X, y)

# Make a prediction for Cyprus
X_new = [[22587]] # Cyprus' GDP per capita
print(model.predict(X_new)) # outputs [[ 5.96242338]]
```

'c' argument looks like a single numeric RGB or RGBA sequence, which should be avoided as value-mapping will have precedence in case its length matches with 'x' & 'y'. Please use a 2-D array with a single row if you really want to specify the same RGB or RGBA value for all points.



```
[[6.25984414]]
```

Step 2: check the dimension of the table

```
[50]: print("The dimension of the table is: ", oecd_bli.shape)
```

The dimension of the table is: (2368, 17)

Step 3: Look at the data

```
[51]: print(oecd_bli.head(5))
```

	LOCATION	Country	INDICATOR	Indicator	MEASURE	\
0	AUS	Australia	JE_LMIS	Labour market insecurity	L	
1	AUT	Austria	JE_LMIS	Labour market insecurity	L	
2	BEL	Belgium	JE_LMIS	Labour market insecurity	L	
3	CAN	Canada	JE_LMIS	Labour market insecurity	L	
4	CZE	Czech Republic	JE_LMIS	Labour market insecurity	L	

	Measure	INEQUALITY	Inequality	Unit	Code	Unit	PowerCode	Code	\
0	Value	TOT	Total	PC	Percentage			0	
1	Value	TOT	Total	PC	Percentage			0	
2	Value	TOT	Total	PC	Percentage			0	
3	Value	TOT	Total	PC	Percentage			0	
4	Value	TOT	Total	PC	Percentage			0	

	PowerCode	Reference Period	Code	Reference Period	Value	Flag Codes	Flags
0	Units		NaN	NaN	5.4	NaN	NaN
1	Units		NaN	NaN	3.5	NaN	NaN
2	Units		NaN	NaN	3.7	NaN	NaN
3	Units		NaN	NaN	6.0	NaN	NaN
4	Units		NaN	NaN	3.1	NaN	NaN

```
[52]: oecd_bli.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2368 entries, 0 to 2367
Data columns (total 17 columns):
LOCATION                2368 non-null object
Country               2368 non-null object
INDICATOR             2368 non-null object
Indicator             2368 non-null object
MEASURE               2368 non-null object
Measure              2368 non-null object
INEQUALITY            2368 non-null object
Inequality            2368 non-null object
Unit Code             2368 non-null object
Unit                 2368 non-null object
PowerCode Code       2368 non-null int64
PowerCode             2368 non-null object
Reference Period Code  0 non-null float64
Reference Period      0 non-null float64
Value                2368 non-null float64
Flag Codes           0 non-null float64
Flags                0 non-null float64
dtypes: float64(5), int64(1), object(11)
memory usage: 314.6+ KB
```

Looking at the results of the “info” method, there are a number of empty columns that can be removed.

There are also some with the same value throughout.

```
[53]: # Remove empty columns
oecd_bli.dropna(axis=1, inplace=True)

# Looks like some other variables have the same value from top to bottom,
# so really don't need them...
print(oecd_bli['MEASURE'].value_counts())
print(oecd_bli['PowerCode Code'].value_counts())
oecd_bli.drop(['MEASURE', 'Measure', 'PowerCode Code'], axis = 1, inplace=True)

# Remove space from Unit Code
oecd_bli.rename(columns={'Unit Code': 'UnitCode'})
```

```
print(oecd_bli.info())
```

```
L      2368
Name: MEASURE, dtype: int64
0      2368
Name: PowerCode Code, dtype: int64
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2368 entries, 0 to 2367
Data columns (total 10 columns):
LOCATION      2368 non-null object
Country      2368 non-null object
INDICATOR    2368 non-null object
Indicator    2368 non-null object
INEQUALITY   2368 non-null object
Inequality   2368 non-null object
Unit Code    2368 non-null object
Unit         2368 non-null object
PowerCode    2368 non-null object
Value        2368 non-null float64
dtypes: float64(1), object(9)
memory usage: 185.1+ KB
None
```

```
[54]: oecd_bli.sample(5)
```

```
[54]:
```

	LOCATION	Country	INDICATOR	Indicator \
748	CAN	Canada	SC_SNTWS	Quality of support network
696	SVK	Slovak Republic	JE_LTUR	Long-term unemployment rate
1645	HUN	Hungary	HS_SFRH	Self-reported health
1786	FIN	Finland	HS_SFRH	Self-reported health
1791	IRL	Ireland	HS_SFRH	Self-reported health

	INEQUALITY	Inequality	Unit	Code	Unit	PowerCode	Value
748	TOT	Total	PC	Percentage	Units	93.00	
696	WMN	Women	PC	Percentage	Units	4.63	
1645	TOT	Total	PC	Percentage	Units	60.00	
1786	LW	Low	PC	Percentage	Units	57.00	
1791	LW	Low	PC	Percentage	Units	72.00	

Using `pandas.pivot_table`, transform dataframe into a more human-friendly format...

```
[55]: df_table = pd.pivot_table(oecd_bli, values='Value', index='Country',
    →columns=['INDICATOR'])

# I also need a 'Country' column. I know this is probably not the right way to
    →go about this...
df_table['country'] = df_table.index.astype('str')
```

```

# Drop this row, it's not a country...
indexNames = df_table[df_table['country'] == 'OECD - Total'].index
df_table.drop(indexNames, inplace=True)

print(df_table.sample(5))

# For reference, create a dictionary of Indicators
print("LIST OF INDICATOR KEYS AND DESCRIPTIONS")
print("=====")
df_indicators = oecd_bli.groupby('INDICATOR')['Indicator'].agg('min')
print(df_indicators.sort_values())

```

INDICATOR	CG_SENG	CG_VOTO	EQ_AIRP	EQ_WATER	ES_EDUA	ES_EDUEX	\
Country							
Ireland	1.3	65.0	7.0	85.000000	82.000000	18.100000	
Netherlands	2.6	82.0	14.0	93.000000	78.333333	18.666667	
Norway	2.2	78.0	5.0	97.666667	82.000000	18.333333	
Hungary	1.2	70.0	19.0	77.000000	84.000000	16.433333	
Slovak Republic	3.0	60.0	21.0	84.666667	91.333333	15.766667	

INDICATOR	ES_STCS	HO_BASE	HO_HISH	HO_NUMR	HS_LEB	HS_SFRH	\
Country							
Ireland	509.8	1.0	20.0	2.1	81.766667	82.8	
Netherlands	509.6	0.1	19.0	1.9	81.600000	75.8	
Norway	504.8	0.0	17.0	2.1	82.466667	77.2	
Hungary	475.0	4.7	19.0	1.2	76.166667	61.0	
Slovak Republic	463.4	1.2	23.0	1.1	77.266667	68.6	

INDICATOR	IW_HADI	IW_HNFW	JE_EMPL	JE_LMIS	JE_LTUR	JE_PEARL	\
Country							
Ireland	25310.0	217130.0	67.333333	8.784	3.203333	47653.0	
Netherlands	29333.0	157824.0	75.666667	4.930	1.976667	52877.0	
Norway	35725.0	228936.0	74.000000	NaN	0.656667	51212.0	
Hungary	NaN	104458.0	68.000000	5.256	1.726667	22576.0	
Slovak Republic	20474.0	NaN	66.000000	21.376	4.773333	24328.0	

INDICATOR	PS_FSAFEN	PS_REPH	SC_SNTWS	SW_LIFS	WL_EWLH	WL_TNOW	\
Country							
Ireland	76.000000	0.666667	95.0	7.050	5.260000	NaN	
Netherlands	82.066667	0.600000	91.0	7.500	0.413333	NaN	
Norway	90.033333	0.400000	94.0	7.620	2.886667	15.563333	
Hungary	56.700000	1.000000	85.5	5.680	2.976667	NaN	
Slovak Republic	63.700000	0.800000	91.5	6.425	4.073333	NaN	

INDICATOR	country
Country	

Ireland Ireland
 Netherlands Netherlands
 Norway Norway
 Hungary Hungary
 Slovak Republic Slovak Republic
 LIST OF INDICATOR KEYS AND DESCRIPTIONS
 =====

INDICATOR
 EQ_AIRP Air pollution
 HO_BASE Dwellings without basic facilities
 ES_EDUA Educational attainment
 WL_EWLH Employees working very long hours
 JE_EMPL Employment rate
 PS_FSAFEN Feeling safe walking alone at night
 PS_REPH Homicide rate
 IW_HADI Household net adjusted disposable income
 IW_HNFW Household net wealth
 HO_HISH Housing expenditure
 JE_LMIS Labour market insecurity
 HS_LEB Life expectancy
 SW_LIFS Life satisfaction
 JE_LTUR Long-term unemployment rate
 JE_PEARN Personal earnings
 SC_SNTWS Quality of support network
 HO_NUMR Rooms per person
 HS_SFRH Self-reported health
 CG_SENG Stakeholder engagement for developing regulations
 ES_STCS Student skills
 WL_TNOW Time devoted to leisure and personal care
 CG_VOTO Voter turnout
 EQ_WATER Water quality
 ES_EDUEX Years in education
 Name: Indicator, dtype: object

[56]: df_table.describe()

[56]:	INDICATOR	CG_SENG	CG_VOTO	EQ_AIRP	EQ_WATER	ES_EDUA	ES_EDUEX	\
	count	38.000000	40.000000	40.000000	40.000000	39.000000	39.000000	
	mean	2.160526	69.57500	13.325000	82.333333	77.717949	17.547863	
	std	0.577291	12.21157	5.770782	10.492977	15.136134	1.412720	
	min	1.200000	47.00000	3.000000	55.333333	37.666667	14.100000	
	25%	1.725000	60.75000	9.750000	74.250000	75.000000	16.550000	
	50%	2.200000	69.50000	14.000000	83.833333	82.000000	17.666667	
	75%	2.575000	79.00000	16.500000	91.083333	87.833333	18.350000	
	max	3.200000	91.00000	28.000000	98.666667	94.000000	20.966667	

	INDICATOR	ES_STCS	HO_BASE	HO_HISH	HO_NUMR	HS_LEB	HS_SFRH	\
	count	39.000000	37.000000	38.000000	37.000000	40.000000	37.000000	

mean	485.707692	5.075676	20.657895	1.632432	79.567500	67.493243
std	33.787972	8.448320	2.528500	0.431441	4.669642	14.331584
min	398.200000	0.000000	15.000000	0.900000	57.500000	33.000000
25%	475.800000	0.300000	19.000000	1.200000	77.916667	60.800000
50%	492.800000	0.900000	21.000000	1.600000	81.366667	70.200000
75%	506.800000	6.700000	22.750000	1.900000	82.366667	76.000000
max	528.800000	37.000000	26.000000	2.600000	84.066667	89.250000

INDICATOR	IW_HADI	IW_HNFW	JE_EMPL	JE_LMIS	JE_LTUR	\
count	29.000000	27.000000	40.000000	33.000000	38.000000	
mean	27807.310345	289780.185185	68.533333	7.706970	2.855789	
std	7055.262661	165673.432787	7.882253	6.234572	3.622899	
min	16275.000000	70160.000000	43.333333	0.662000	0.050000	
25%	21453.000000	180100.000000	65.833333	4.392000	1.011667	
50%	29333.000000	259667.000000	69.666667	5.396000	1.776667	
75%	31304.000000	379777.000000	74.000000	8.784000	3.196667	
max	45284.000000	769053.000000	85.666667	29.200000	16.643333	

INDICATOR	JE_PEARL	PS_FSAFEN	PS_REPH	SC_SNTWS	SW_LIFS	\
count	35.000000	40.000000	40.000000	40.000000	40.000000	
mean	39817.514286	68.463333	3.481667	90.193333	6.577208	
std	13108.329748	13.960934	6.459861	4.384954	0.762724	
min	15314.000000	35.866667	0.166667	78.333333	4.700000	
25%	25971.500000	60.108333	0.600000	88.300000	5.938333	
50%	40863.000000	70.483333	0.950000	91.350000	6.510000	
75%	49400.500000	78.500000	2.166667	93.062500	7.243750	
max	63062.000000	90.033333	27.000000	98.000000	7.660000	

INDICATOR	WL_EWLH	WL_TNOW
count	38.000000	22.000000
mean	7.789649	15.048939
std	7.585983	0.672978
min	0.140000	13.826667
25%	3.150833	14.560833
50%	4.981667	14.885000
75%	10.571667	15.600833
max	31.043333	16.336667

```
[57]: corr_matrix = df_table.corr()
      corr_matrix["SW_LIFS"].sort_values(ascending=False)
```

```
[57]: INDICATOR
      SW_LIFS      1.000000
      JE_PEARL    0.731418
      IW_HADI     0.713008
      EQ_WATER    0.682587
      JE_EMPL     0.678344
      SC_SNTWS    0.667896
```



```

HS_SFRH      0.656817
PS_FSAFEN    0.600163
HO_NUMR      0.597502
HS_LEB       0.568044
CG_VOTO      0.368598
ES_EDUEX     0.324655
ES_EDUA      0.293395
IW_HNFW      0.292887
HO_HISH      0.286334
WL_TNOW      0.199424
ES_STCS      0.197223
CG_SENG      0.180861
WL_EWLH      -0.195136
PS_REPH      -0.259378
JE_LMIS      -0.452874
HO_BASE      -0.528167
EQ_AIRP      -0.551376
JE_LTUR      -0.567002
Name: SW_LIFS, dtype: float64

```

Step 4: Think about some questions that might help you predict what indicators most influence the Life Satisfaction score: The central point of this dataset is the so-called, “**Life Satisfaction Index**”. In other words, do indicators in the categories of housing, income, jobs, community, education, environment, civic engagement, health, etc. really lead to a better, more satisfied life? Let’s focus on a few high-level categories to see how the indicators correlate with the LSI...

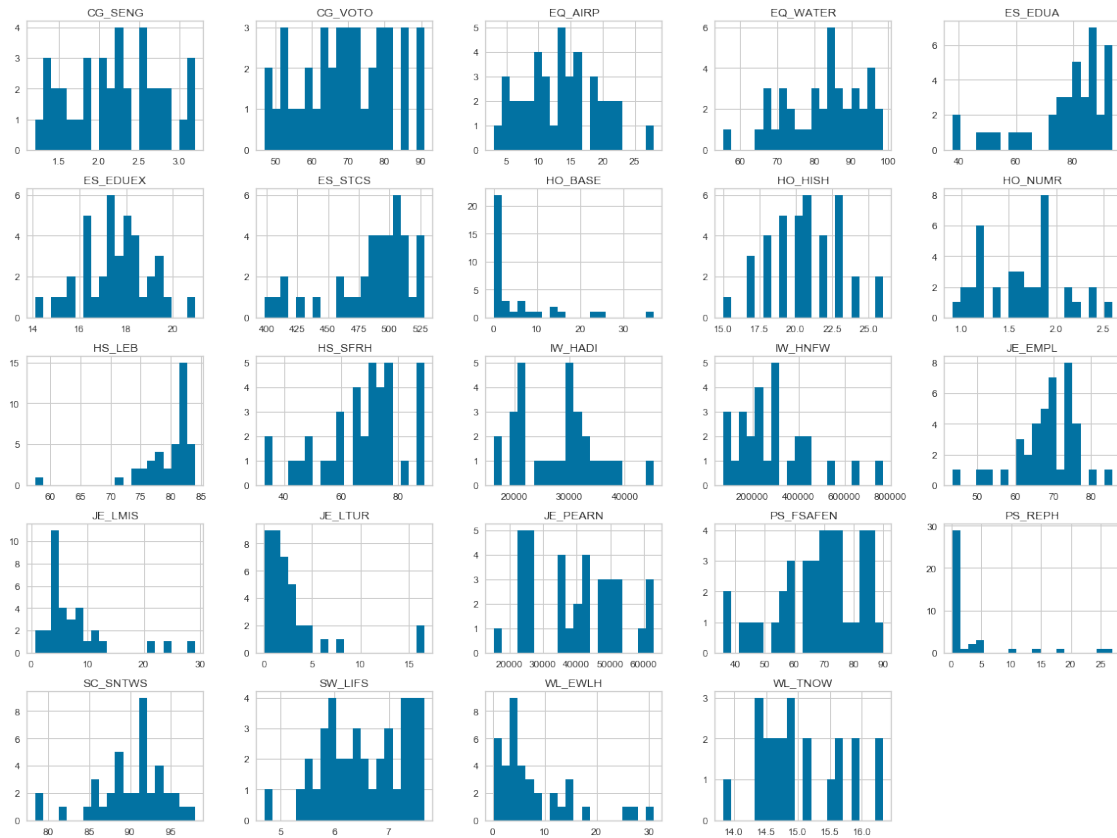
- **Wealth** Net Wealth, Labor Market Insecurity, Employment rate
- **Environment** Air pollution, Homicide rate, Water quality
- **Health** Life expectancy, Self-reported health, Long work hours

My observations are based on the **graph analysis** below.

First, there does not seem to be any noticeable normal distribution amongst any of the indicators, although some—such as HS_LEB (Life Expectancy) exhibit *normal-ish* distribution on a skewed scale.

- **Wealth** - somewhat surprisingly, Net Wealth does not appear to be as important as labor market security and the employment rate. Having said that, removing the data points above \$500,000 might tell a different story.
- **Environment** - Air and water quality seem to factor higher than the homicide rate, which shows almost no effect on the LSI.
- **Health** - Life expectancy seems like an obvious factor, but I was also satisfied to see long work hours affect the index as well.

```
[58]: df_table.hist( bins = 20, figsize =( 20,15))
      plt.show()
```



```
[59]: results = smf.ols('SW_LIFS ~ IW_HNFW + JE_LMIS + JE_EMPL', data=df_table).fit()
print("Money: Net Wealth, Labor Market Insecurity, Employment rate")
print(results.summary())
```

Money: Net Wealth, Labor Market Insecurity, Employment rate

OLS Regression Results

```
=====
Dep. Variable:          SW_LIFS    R-squared:                0.327
Model:                  OLS        Adj. R-squared:           0.235
Method:                 Least Squares    F-statistic:           3.564
Date:                   Sat, 01 Feb 2020    Prob (F-statistic):    0.0306
Time:                   13:14:13    Log-Likelihood:       -22.710
No. Observations:      26    AIC:                   53.42
Df Residuals:          22    BIC:                   58.45
Df Model:               3
Covariance Type:       nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	2.7241	2.217	1.229	0.232	-1.873	7.321
IW_HNFW	1.13e-06	7.6e-07	1.486	0.151	-4.47e-07	2.71e-06

JE_LMIS	-0.0085	0.028	-0.304	0.764	-0.067	0.050
JE_EMPL	0.0523	0.030	1.763	0.092	-0.009	0.114
=====						
Omnibus:		0.953	Durbin-Watson:			1.785
Prob(Omnibus):		0.621	Jarque-Bera (JB):			0.831
Skew:		-0.167	Prob(JB):			0.660
Kurtosis:		2.190	Cond. No.			6.02e+06
=====						


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The condition number is large, 6.02e+06. This might indicate that there are strong multicollinearity or other numerical problems.

```
[60]: attributes = ['SW_LIFS', 'IW_HNFW', 'JE_LMIS', 'JE_EMPL']
scatter_matrix(df_table[attributes], alpha=1.0, figsize=(12, 8))
```

```
[60]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E299AD30>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E2982700>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E29FAEE0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E2A11730>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E2A24F70>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E2A3C7F0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E2A44700>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E2A5FFA0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E2AA7EE0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E4F1C970>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E4F54FA0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E4F7AA30>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E4FB22B0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E4FDBAF0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E56A1370>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E56CABB0>]],
dtype=object)
```



```
[61]: (ggplot(df_table, aes(x='IW_HNFW', y='SW_LIFS', color='JE_LMIS',
→size='JE_EMPL')) +
      geom_point(alpha=0.5) +
      scale_x_log10() +
      geom_text(aes(x='IW_HNFW', y='SW_LIFS', label='country'),
                color="black",
                size=8,
                data=df_table) +
      theme(figure_size = (12.0, 8.0)) +
      labs(title="Life Satisfaction versus Wealth Factors",x="Net_
→Wealth",y="Life Satisfaction Index",size="Employment Rate",color="Labor_
→Market Insecurity")
)
```




[61]: <ggplot: (156738463305)>

```
[62]: import yellowbrick
from yellowbrick.features import Rank2D
from yellowbrick.features import ParallelCoordinates
from yellowbrick.style import set_palette

#set up the figure size
plt.rcParams['figure.figsize'] = (15, 7)
num_features = ['SW_LIFS', 'IW_HNFW', 'JE_LMIS', 'JE_EMPL']
# extract the numpy arrays from the data frame
X = df_table[num_features].as_matrix()

# instantiate the visualizer with the Covariance ranking algorithm
visualizer = Rank2D(features=num_features, algorithm='pearson')
visualizer.fit(X) # Fit the data to the visualizer
visualizer.transform(X) # Transform the data

plt.show()
```



```
[63]: results = smf.ols('SW_LIFS ~ EQ_AIRP + PS_REPH + EQ_WATER', data=df_table).fit()
print("Environment: Air pollution, Homicide rate, Water quality")
print(results.summary())
```

Environment: Air pollution, Homicide rate, Water quality

OLS Regression Results

```
=====
Dep. Variable:          SW_LIFS    R-squared:                0.512
Model:                  OLS        Adj. R-squared:           0.471
Method:                 Least Squares    F-statistic:             12.57
Date:                   Sat, 01 Feb 2020    Prob (F-statistic):       8.92e-06
Time:                   13:14:15          Log-Likelihood:          -31.080
No. Observations:       40              AIC:                   70.16
Df Residuals:           36              BIC:                   76.92
Df Model:                3
Covariance Type:        nonrobust
=====
```

```
=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
```

Intercept	3.8394	1.218	3.152	0.003	1.369	6.310
EQ_AIRP	-0.0343	0.020	-1.751	0.088	-0.074	0.005
PS_REPH	-0.0015	0.017	-0.093	0.926	-0.035	0.032
EQ_WATER	0.0389	0.012	3.193	0.003	0.014	0.064

```
=====
Omnibus:                3.892    Durbin-Watson:                1.767
Prob(Omnibus):          0.143    Jarque-Bera (JB):        2.842
Skew:                   -0.368    Prob(JB):                0.241
Kurtosis:               4.079    Cond. No.                1.17e+03
=====
```


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The condition number is large, 1.17e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
[64]: attributes = ['SW_LIFS', 'EQ_AIRP', 'PS_REPH', 'EQ_WATER']
      scatter_matrix(df_table[attributes], alpha=1.0, figsize=(12, 8))
```

```
[64]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E4E78130>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E27BF460>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E28BF970>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E29575B0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E5229E50>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247CBA952E0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E2C34730>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E51012B0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E2818040>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E28DE820>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E4C29D00>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E4A53FA0>],
[<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E4C1B400>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E49AFA30>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E4C4DAF0>,
<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E285DB80>]],
dtype=object)
```



```
[65]: (ggplot(df_table, aes(x='EQ_AIRP', y='SW_LIFS', color='PS_REPH',
→size='EQ_WATER')) +
      geom_point(alpha = 0.5) +
      geom_text(aes(x='EQ_AIRP', y='SW_LIFS', label='country'),
                color="black",
                size=8,
                data=df_table) +
      theme(figure_size = (12.0, 8.0)) +
      labs(title="Life Satisfaction versus Environment",x="Air_
→Quality",y="Life Satisfaction Index",size="Water Quality",color="Homicide_
→Rate")
    )
```





[65]: <ggplot: (156737894614)>

```
[66]: #set up the figure size
plt.rcParams['figure.figsize'] = (15, 7)
num_features = ['SW_LIFS', 'EQ_AIRP', 'PS_REPH', 'EQ_WATER']
# extract the numpy arrays from the data frame
X = df_table[num_features].as_matrix()

# instantiate the visualizer with the Covariance ranking algorithm
visualizer = Rank2D(features=num_features, algorithm='pearson')
visualizer.fit(X) # Fit the data to the visualizer
visualizer.transform(X) # Transform the data

plt.show()
```



```
[67]: results = smf.ols('SW_LIFS ~ HS_LEB + HS_SFRH + WL_EWLH', data=df_table).fit()
print("Health: Life expectancy, Self-reported health, Long work hours")
print(results.summary())
```

Health: Life expectancy, Self-reported health, Long work hours

OLS Regression Results

```
=====
Dep. Variable:          SW_LIFS    R-squared:                0.484
Model:                  OLS        Adj. R-squared:           0.434
Method:                 Least Squares    F-statistic:             9.707
Date:                   Sat, 01 Feb 2020    Prob (F-statistic):       0.000114
Time:                   13:14:16          Log-Likelihood:          -26.111
No. Observations:       35              AIC:                    60.22
Df Residuals:           31              BIC:                    66.44
Df Model:                3
Covariance Type:        nonrobust
=====
```

```
=====
               coef      std err          t      P>|t|      [0.025      0.975]
-----
```

Intercept	2.3807	3.176	0.750	0.459	-4.096	8.858
HS_LEB	0.0227	0.046	0.499	0.621	-0.070	0.116
HS_SFRH	0.0382	0.011	3.403	0.002	0.015	0.061
WL_EWLH	-0.0250	0.014	-1.737	0.092	-0.054	0.004

Omnibus:	8.285	Durbin-Watson:	2.096
Prob(Omnibus):	0.016	Jarque-Bera (JB):	6.935
Skew:	-0.902	Prob(JB):	0.0312
Kurtosis:	4.227	Cond. No.	3.70e+03


Warnings:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The condition number is large, 3.7e+03. This might indicate that there are strong multicollinearity or other numerical problems.

```
[68]: attributes = ['SW_LIFS', 'HS_LEB', 'HS_SFRH', 'WL_EWLH']
      scatter_matrix(df_table[attributes], alpha=1.0, figsize=(12, 8))
```

```
[68]: array([[<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E542DC10>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E5047460>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E53D2C40>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E510CD90>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E5171F70>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E527A6D0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E5469520>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247DF9CEA00>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x00000247E5302490>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E538E910>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E5227D90>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E51CC310>],
      [<matplotlib.axes._subplots.AxesSubplot object at 0x00000247DFD57B50>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E52D83D0>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E5339C10>,
      <matplotlib.axes._subplots.AxesSubplot object at 0x00000247E5520490>]],
      dtype=object)
```



```
[69]: (ggplot(df_table[df_table['HS_LEB'] > 70], aes(x='HS_LEB', y='SW_LIFS',
→color='WL_EWLH', size='HS_SFRH')) +
  geom_point(alpha = 0.8) +
  scale_x_continuous(limits=[70,85]) +
  geom_text(aes(x='HS_LEB', y='SW_LIFS', label='country'),
    color="black",
    size=8,
    data=df_table) +
  theme(figure_size = (12.0, 8.0)) +
  labs(title="Life Satisfaction versus Health Factors",x="Life_
→Expectancy",y="Life Satisfaction Index",size="Self-reported_
→Health",color="Long Work Hours")
)
```




[69]: <ggplot: (156737931379)>

```
[70]: #set up the figure size
plt.rcParams['figure.figsize'] = (15, 7)
num_features = ['SW_LIFS', 'HS_LEB', 'HS_SFRH', 'WL_EWLH']
# extract the numpy arrays from the data frame
X = df_table[num_features].as_matrix()

# instantiate the visualizer with the Covariance ranking algorithm
visualizer = Rank2D(features=num_features, algorithm='pearson')
visualizer.fit(X) # Fit the data to the visualizer
visualizer.transform(X) # Transform the data

plt.show()
```



Part 2: Dimensionality and Feature Reduction

```
[71]: # What kind of data are we dealing with?
df_table.describe()
```

```
[71]: INDICATOR    CG_SENG    CG_VOTO    EQ_AIRP    EQ_WATER    ES_EDUA    ES_EDUEX  \
count      38.000000   40.00000   40.000000   40.000000   39.000000   39.000000
mean         2.160526   69.57500   13.325000   82.333333   77.717949   17.547863
std          0.577291   12.21157    5.770782   10.492977   15.136134    1.412720
min          1.200000   47.00000    3.000000   55.333333   37.666667   14.100000
25%          1.725000   60.75000    9.750000   74.250000   75.000000   16.550000
50%          2.200000   69.50000   14.000000   83.833333   82.000000   17.666667
75%          2.575000   79.00000   16.500000   91.083333   87.833333   18.350000
max          3.200000   91.00000   28.000000   98.666667   94.000000   20.966667

INDICATOR    ES_STCS    HO_BASE    HO_HISH    HO_NUMR    HS_LEB    HS_SFRH  \
count      39.000000   37.000000   38.000000   37.000000   40.000000   37.000000
```

mean	485.707692	5.075676	20.657895	1.632432	79.567500	67.493243
std	33.787972	8.448320	2.528500	0.431441	4.669642	14.331584
min	398.200000	0.000000	15.000000	0.900000	57.500000	33.000000
25%	475.800000	0.300000	19.000000	1.200000	77.916667	60.800000
50%	492.800000	0.900000	21.000000	1.600000	81.366667	70.200000
75%	506.800000	6.700000	22.750000	1.900000	82.366667	76.000000
max	528.800000	37.000000	26.000000	2.600000	84.066667	89.250000

INDICATOR	IW_HADI	IW_HNFW	JE_EMPL	JE_LMIS	JE_LTUR	\
count	29.000000	27.000000	40.000000	33.000000	38.000000	
mean	27807.310345	289780.185185	68.533333	7.706970	2.855789	
std	7055.262661	165673.432787	7.882253	6.234572	3.622899	
min	16275.000000	70160.000000	43.333333	0.662000	0.050000	
25%	21453.000000	180100.000000	65.833333	4.392000	1.011667	
50%	29333.000000	259667.000000	69.666667	5.396000	1.776667	
75%	31304.000000	379777.000000	74.000000	8.784000	3.196667	
max	45284.000000	769053.000000	85.666667	29.200000	16.643333	

INDICATOR	JE_PEARL	PS_FSAFEN	PS_REPH	SC_SNTWS	SW_LIFS	\
count	35.000000	40.000000	40.000000	40.000000	40.000000	
mean	39817.514286	68.463333	3.481667	90.193333	6.577208	
std	13108.329748	13.960934	6.459861	4.384954	0.762724	
min	15314.000000	35.866667	0.166667	78.333333	4.700000	
25%	25971.500000	60.108333	0.600000	88.300000	5.938333	
50%	40863.000000	70.483333	0.950000	91.350000	6.510000	
75%	49400.500000	78.500000	2.166667	93.062500	7.243750	
max	63062.000000	90.033333	27.000000	98.000000	7.660000	

INDICATOR	WL_EWLH	WL_TNOW
count	38.000000	22.000000
mean	7.789649	15.048939
std	7.585983	0.672978
min	0.140000	13.826667
25%	3.150833	14.560833
50%	4.981667	14.885000
75%	10.571667	15.600833
max	31.043333	16.336667

Proposed Steps for Dimensionality and Feature Reduction

- The *variance* between features varies greatly, For example, CG_SENS ranges from 1.2 to 3.2 while SC_SNTWS ranges from 4.38 to 98. I will start by setting a variance threshold and eliminating features below the threshold. In theory, this should eliminate variables with low variance, which likely will not contribute greatly to the model.
- Due to the quantitative nature of these variables, there exists wide discrepancies of units. For instance, Household Net Wealth can be in the thousands of US dollars, while Employment Rate is a percentage. I will apply both a simple min-max rescaler to normalize all variables, and a standard scaler to standardize the variables.

- I suspect that variables such as air pollution and water quality probably are highly correlated. I will apply a correlation matrix and will consider dropping one of the correlated features.
- Lastly, I'd like to automatically select the best features to keep by leveraging scikit-learn's recursive feature elimination functionality.

VarianceThreshold

```
[72]: # Make a features DataFrame by dropping the target variable--SW_LIFS--and
      ↪ assigning that to a target Series,
features = df_table.drop(['country', 'SW_LIFS'], axis=1)
target = df_table['SW_LIFS']
```

```
[73]: # Apply a Variance Threshold to remove those features with low variance
      from sklearn.feature_selection import VarianceThreshold

      # Make a features dataset by dropping the target variable--SW_LIFS
features = df_table.drop(['country', 'SW_LIFS'], axis=1)
target = df_table['SW_LIFS']

      # Create thresholder
thresholder = VarianceThreshold(threshold=5.0)

      # Create high variance feature matrix
features_high_variance = thresholder.fit_transform(features)

features_vt = features.loc[:, thresholder.get_support()]
features_vt.head()
```

```
[73]: INDICATOR  CG_VOTO  EQ_AIRP  EQ_WATER  ES_EDUA  ES_STCS  HO_BASE  HO_HISH  \
Country
Australia    91.0      5.0  92.666667  81.000000  411.2      NaN      20.0
Austria      80.0     16.0  92.000000  85.000000  492.8      0.9     21.0
Belgium      89.0     15.0  83.666667  77.000000  503.8      1.9     21.0
Brazil       79.0     10.0  73.000000  49.000000  398.2      6.7      NaN
Canada       68.0      7.0  91.000000  91.333333  523.2      0.2     22.0
```

```
INDICATOR    HS_LEB  HS_SFRH  IW_HADI  IW_HNFW  JE_EMPL  JE_LMIS  \
Country
Australia  82.500000   87.25  32759.0  427064.0  73.000000   5.922
Austria    81.700000   70.60  33541.0  308325.0  72.000000   4.076
Belgium    81.500000   73.60  30364.0  386006.0  63.333333   4.052
Brazil     74.766667    NaN     NaN     NaN    61.000000    NaN
Canada     81.866667   87.80  30854.0  423849.0  73.333333   7.048
```

```
INDICATOR    JE_LTUR  JE_PEARL  PS_FSAFEN  PS_REPH  SC_SNTWS  WL_EWLH
Country
Australia    1.306667  49126.0  64.133333   1.100000   95.25  12.840000
Austria      1.830000  50349.0  80.700000   0.466667   92.00   6.590000
Belgium      3.533333  49675.0  70.266667   1.033333   92.00   4.703333
```


Brazil	NaN	NaN	35.866667	27.000000	89.25	7.006667
Canada	0.763333	47622.0	82.500000	1.266667	93.25	3.673333

Not terribly surprising to see which features the Variance Thesholder dropped:

Stakeholder engagement for developing regulations

Years in education

Rooms per person

Time devoted to leisure and personal care

```
[74]: from sklearn import preprocessing
```

```
# Standardizing
```

```
std_scale = preprocessing.StandardScaler().fit(features_vt)
```

```
df_std = std_scale.transform(features_vt)
```

```
standardized_df = pd.DataFrame(df_std, index=features_vt.index,
    ↳columns=features_vt.columns)
```

```
#Min-max scaling (normalization)
```

```
minmax_scale = preprocessing.MinMaxScaler().fit(features_vt)
```

```
df_minmax = minmax_scale.transform(features_vt)
```

```
scaled_features_df = pd.DataFrame(df_minmax, index=features_vt.index,
    ↳columns=features_vt.columns)
```

```
[75]: def plot():
```

```
    plt.figure(figsize=(8,6))
```

```
    plt.scatter(df_table['JE_EMPL'], df_table['IW_HNFW'],
        color='green', label='original scale', alpha=0.5)
```

```
    plt.scatter(df_std[:,4], df_std[:,13], color='red',
        label='standardized [mu=0,sigma=1]', alpha=0.3)
```

```
    plt.scatter(df_minmax[:,4], df_minmax[:,13],
        color='blue', label='min-max scaled [min=0, max=1]', alpha=0.3)
```

```
    plt.title('Employment Rate and Household Income')
```

```
    plt.xlabel('Empl Rate')
```

```
    plt.ylabel('HH Income')
```

```
    plt.legend(loc='upper left')
```

```
    plt.xscale('symlog')
```

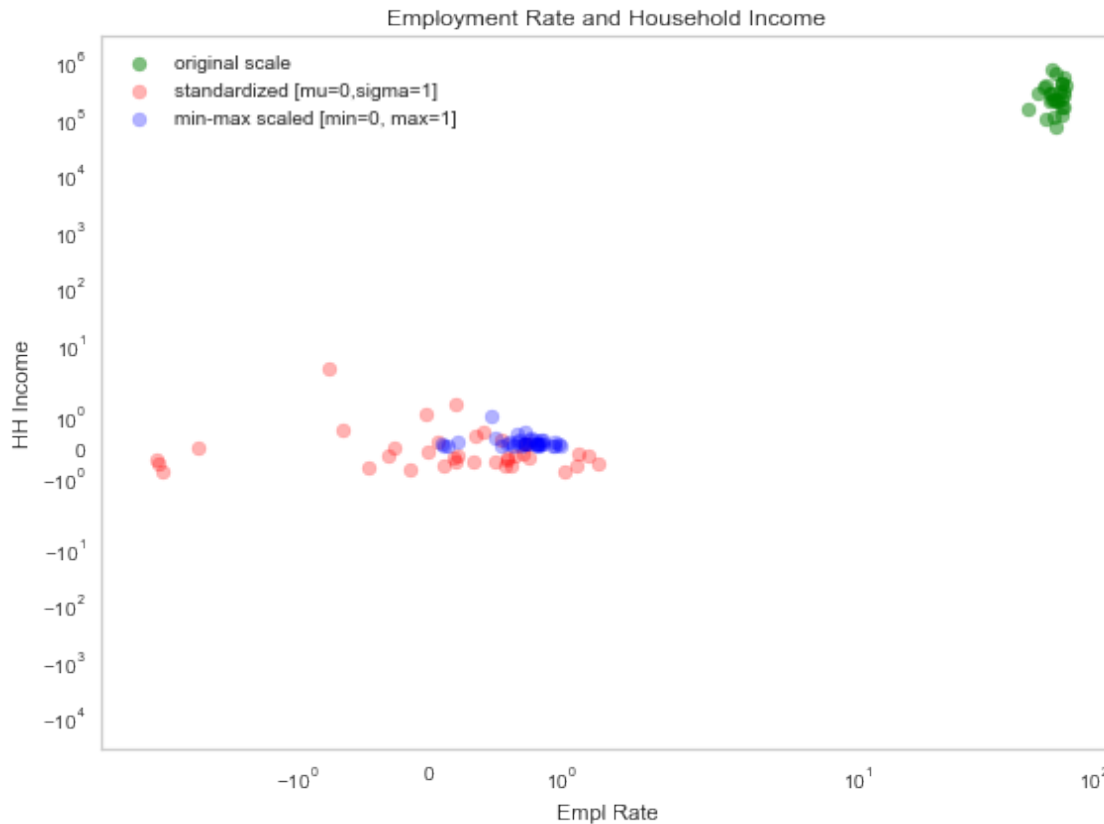
```
    plt.yscale('symlog')
```

```
    plt.grid()
```

```
    plt.tight_layout()
```

```
plot()
```

```
plt.show()
```



At this point I made a more or less arbitrary decision to go with the standardized feature set over the min max scaled...

```
[82]: # Create correlation matrix
corr_matrix = standardized_df.corr().abs()

# Select upper triangle of correlation matrix
upper = corr_matrix.where(np.triu(np.ones(corr_matrix.shape),
                                   k=1).astype(np.bool))
upper.fillna(0, inplace=True)

# Find index of feature columns with correlation greater than 0.95
to_drop = [column for column in upper.columns if any(upper[column] > 0.95)]

# Drop features
standardized_df.drop(standardized_df.columns[to_drop], axis=1, inplace=True)
features = standardized_df
print(features.shape)
features.sample(5)
```

(40, 19)

```
[82]: INDICATOR      CG_VOTO  EQ_AIRP  EQ_WATER  ES_EDUA  ES_STCS  HO_BASE  \
Country
Czech Republic -0.711149  1.171425  0.418236  1.067459  0.194660 -0.525079
Mexico         -0.545283  0.469447 -1.415567 -2.680661 -2.090062  2.450908
France         0.449910 -0.057036 -0.128688  0.041188  0.350573 -0.549079
New Zealand    0.864574 -1.460990  0.643439  0.063498  0.614426      NaN
Lithuania     -1.540476  0.118459 -0.160860  1.000529 -0.273078  1.022914


INDICATOR      HO_HISH  HS_LEB  HS_SFRH  IW_HADI  IW_HNFW  JE_EMPL  \
Country
Czech Republic  1.339516 -0.101390 -0.473470 -0.916590      NaN  0.659549
Mexico         -0.263684 -0.903836 -0.129209      NaN      NaN -0.882253
France         0.137116  0.607076 -0.077334  0.504387 -0.056141 -0.411147
New Zealand    2.141117  0.462491  1.539039      NaN  0.607307  1.087827
Lithuania     -0.664485 -1.033963 -1.548694 -0.886731      NaN  0.231270

INDICATOR      JE_LMIS  JE_LTUR  JE_PEARL  PS_FSAFEN  PS_REPH  SC_SNTWS  \
Country
Czech Republic -0.081274 -0.503264 -1.118099  0.295242 -0.467448  0.359524
Mexico         -0.376417 -0.780194 -1.896599 -1.922093  2.375390 -1.984698
France         0.109625  0.319134  0.304766  0.162250 -0.472674 -0.275610
New Zealand   -0.489784 -0.592776  0.017453 -0.159348 -0.342029  1.398835
Lithuania      NaN -0.044511 -1.202079 -0.887176 -0.007577 -0.414184

INDICATOR      WL_EWLH
Country
Czech Republic -0.306323
Mexico         2.603747
France         -0.017765
New Zealand    0.968140
Lithuania     -0.967601
```

```
[83]: import seaborn as sns
      # print(corr_matrix)
      sns.heatmap(corr_matrix, annot=True)
```

```
[83]: <matplotlib.axes._subplots.AxesSubplot at 0x247e2c8a0d0>
```



There are no features with a correlation of $> 95\%$. If I drop the threshold to 90% there is one: JE_PEARL and IW_HADI (Personal Income and Household disposable income) Although not surprising as they are similar measurements, I will leave them both in.

```
[86]: # Apply Recursive Feature Elimination

# Load libraries
import warnings
from sklearn.feature_selection import RFECV
from sklearn.ensemble import RandomForestRegressor

# Suppress an annoying but harmless warning
warnings.filterwarnings(action="ignore", module="scipy",
                        message="^internal gelsd")

# Need to fix the NaN's before we can successfully apply a model
features = features.fillna(features.median())
randomforest = RandomForestRegressor(random_state=0, n_jobs=-1)

# Recursively eliminate features
rfecv = RFECV(estimator=randomforest, step=1, scoring="neg_mean_squared_error")
rfecv.fit(features, target)
rfecv.transform(features)
```

Number of features we should keep: 15

```
[97]: # Once we have conducted RFE, we can see the number of features we should keep:
# Number of best features
print("Number of features we should keep: {}".format(rfecv.n_features_))
print()

# We can also see which of those features we should keep:
# Which categories are best
print(rfecv.support_)

# Rank features best (1) to worst
print(rfecv.ranking_)

df_rfecv = pd.DataFrame(features.columns, rfecv.ranking_)
print(df_rfecv.sort_values)
```

Number of features we should keep: 15

```
[False False  True  True  True  True  True  True  True  True False  True
  True  True  True  True False  True  True]
[5 3 1 1 1 1 1 1 1 1 4 1 1 1 1 1 2 1 1]
<bound method DataFrame.sort_values of      INDICATOR
5      CG_VOTO
3      EQ_AIRP
1      EQ_WATER
1      ES_EDUA
1      ES_STCS
1      HO_BASE
1      HO_HISH
1      HS_LEB
1      HS_SFRH
1      IW_HADI
4      IW_HNFW
1      JE_EMPL
1      JE_LMIS
1      JE_LTUR
1      JE_PEARL
1      PS_FSAFEN
2      PS_REPH
1      SC_SNTWS
1      WL_EWLH>
```

I have mixed feelings about these results, so I opted to drop the two with the lowest ranking: CG_VOTO (Voter Turnout) and IW_HNFW (Household Net Wealth).

```
[101]: features.drop(['CG_VOTO', 'IW_HNFW'], axis=1, inplace=True)
features.sample(3)
```

```
[101]: INDICATOR      EQ_AIRP  EQ_WATER  ES_EDUA  ES_STCS  HO_BASE  HO_HISH  \
Country
```

Slovak Republic	1.346919	0.225204	0.911288	-0.668857	-0.465079	0.938716
Poland	1.522413	-0.032172	0.955908	0.590439	-0.249080	0.537916
Luxembourg	-0.232530	0.160860	-0.070363	-0.009226	-0.549079	0.137116

INDICATOR	HS_LEB	HS_SFRH	IW_HADI	JE_EMPL	JE_LMIS	JE_LTUR	\
Country							
Slovak Republic	-0.498999	0.078290	-1.057808	-0.325491	2.226451	0.536389	
Poland	-0.347184	-0.530060	-1.153011	-0.325491	-0.117759	-0.375522	
Luxembourg	0.693827	0.134881	1.652593	-0.282664	-0.974849	-0.143348	

INDICATOR	JE_PEARL	PS_FSAFEN	PS_REPH	SC_SNTWS	WL_EWLH
Country					
Slovak Republic	-1.198906	-0.345537	-0.420416	0.301785	-0.496468
Poland	-0.988529	-0.072299	-0.430868	-0.968483	-0.262683
Luxembourg	1.799149	0.537045	-0.456997	0.648222	-0.540999

0.1.2 Part 3: Model Evaluation and Selection

Model Selection

I've decided to start with a **Random Forest Regression** model, leveraging scikit-learn's `RandomForestRegressor`. This model takes advantage of the *ensemble learning method*, where multiple decision trees are trained, although each tree only receives a "bootstrapped" sample of observations. [1]

In order to "dial in" my model, I plan to leverage the `feature_importances_` of the `RandomForestClassifier` model, as well as `SelectFromModel` method from scikit-learn's `feature_selection` library.

Model Evaluation

This topic can get a little overwhelming, as, from the reading I've done, there appears to be a bewildering number of approaches to model evaluation. Given the constraints of one semester's worth of time, I'm going to focus on the following:

- Calculate the **Mean Squared Error**
- Calculate the **Coefficient of Determination (*R squared*)**

References

[1] Albon, Chris. Machine Learning with Python Cookbook: Practical Solutions from Preprocessing to Deep Learning (p. 238). O'Reilly Media. Kindle Edition.

Model Selection

```
[102]: # Apply a Random Forest Regressor Model and train (fit) it.
from sklearn.ensemble import RandomForestRegressor

# Create random forest classifier object
randomforest = RandomForestRegressor(random_state=0, n_jobs=-1)

# Train model
```

```
model = randomforest.fit(features, target)
```

```
[103]: # Calculate feature importances
importances = model.feature_importances_

# Sort feature importances in descending order
indices = np.argsort(importances)[::-1]

# Rearrange feature names so they match the sorted feature importances
names = [features.columns[i] for i in indices]
```


```
[104]: # Create plot
plt.figure()

# Create plot title
plt.title("Feature Importance")

# Add bars
plt.bar(range(features.shape[1]), importances[indices])

# Add feature names as x-axis labels
plt.xticks(range(features.shape[1]), names, rotation=90)

# Show plot
plt.show()
```



Based on the results above, I'm dropping all but the highest rated columns above 0.5.

Interesting observation. I've been resisting dropping EQ_AIRP–Air Quality–because I just can't believe it's not important. But many of the calculations I've run have suggested eliminating it.

```
[107]: new_features = ['JE_PEARL', 'EQ_WATER', 'SC_SNTWS', 'HS_SFRH', 'JE_EMPL', 'HO_BASE']
features_final = features[new_features]

# Re-train the model on the new feature set
model = randomforest.fit(features_final, target)
```

```
[ ]:
```

SelectFromModel is returning dubious results (I'm sure it's something I'm doing wrong) so dropping this for now...

```
[ ]:
```

```
[ ]: # from sklearn.feature_selection import SelectFromModel

# # Create object that selects features with importance greater
# # than or equal to a threshold
# selector = SelectFromModel(randomforest, threshold=0.3)

# # Feature new feature matrix using selector
# features_important = selector.fit_transform(features, target)

# # Sort feature importances in descending order
# indices = np.argsort(features_important)[:, :-1]

# # Rearrange feature names so they match the sorted feature importances
# names = [features.columns for i in indices]

# # Create plot
# plt.figure()

# # Create plot title
# plt.title("Feature Importance")

# # Add bars
# plt.bar(range(features.shape[1]), features_important[indices])

# # Add feature names as x-axis labels
# plt.xticks(range(features.shape[1]), names, rotation=90)

# # Show plot
# plt.show()

# Train random forest using most important features
# model = randomforest.fit(features_important, target)
```

Model Evaluation


```
[125]: # Cross-validate the random forest regression using (negative) MSE, RMSE and
      ↪ R-squared
from sklearn.model_selection import cross_val_score
scorers = ['neg_mean_absolute_error', 'neg_root_mean_squared_error', 'r2']

for scorer in scorers:
    cvs = cross_val_score(randomforest, features, target, scoring=scorer)
    print("{}\t\t{}".format(scorer, cvs))
```

```
neg_mean_absolute_error      [-0.2856625 -0.46934375 -0.3044625 -0.28794583
-0.47037917]
neg_root_mean_squared_error  [-0.33881842 -0.5923367 -0.49264287
-0.35860543 -0.66774549]
r2                            [0.36472877 0.56892412 0.16964655 0.76679536 0.46711817]
```

References http://sebastianraschka.com/Articles/2014_about_feature_scaling.html#standardization-and-min-max-scaling

Albon, Chris. Machine Learning with Python Cookbook: Practical Solutions from Preprocessing to Deep Learning . O'Reilly Media. Kindle Edition.